

tidyr, dplyr and ggplot2

Ken Butler
Lecturer (Statistics), UTSC

Greater Toronto Area R Users Group

November 4, 2015

Does this give you a headache?

DrugA	DrugB	DrugC
4	6	6
5	8	7
4	4	6
3	5	6
2	4	7
4	6	5
3	5	6
4	8	5
4	6	5

- ▶ 27 migraine-suffering subjects, randomly chosen to receive *one* pain-killing drug
- ▶ On next migraine episode, take chosen drug, report pain 30 minutes after (1=no pain, 10=extreme pain)
- ▶ How do drugs compare?

The problem

- ▶ All numbers in data are pain scores, even though in 3 columns.
- ▶ “Wide format”.
- ▶ For eg. boxplots, ANOVA, want 2 columns:
 - ▶ one containing *all* scores
 - ▶ one identifying drug.
- ▶ How to get that?

Tidy data (Wickham)

- ▶ Every value belongs to a *variable* and an *observation*.
- ▶ Variables in columns.
- ▶ Observations in rows.
- ▶ If this is done, data called “tidy”, ready for further analysis.
- ▶ If not, have “untidy” data, needs tidying.
- ▶ “Tidy” depends (somewhat) on kind of analysis you want to do.

Our data

```
> migraine=read.table("migraine.txt",header=T)
> migraine
```

	DrugA	DrugB	DrugC
1	4	6	6
2	5	8	7
3	4	4	6
4	3	5	6
5	2	4	7
6	4	6	5
7	3	5	6
8	4	8	5
9	4	6	5

- ▶ 3 columns all one variable. Not tidy!

gather: combining columns

- ▶ Combine columns that all measure same thing.
- ▶ Input: data frame, what makes columns different, what makes them same, columns to combine:

```
> library(tidyr)
```

```
> migraine2=gather(migraine,drug,pain,DrugA:DrugC)
```


The result

> migraine2			8 DrugA 4	18 DrugB 6
			9 DrugA 4	19 DrugC 6
drug pain			10 DrugB 6	20 DrugC 7
1	DrugA	4	11 DrugB 8	21 DrugC 6
2	DrugA	5	12 DrugB 4	22 DrugC 6
3	DrugA	4	13 DrugB 5	23 DrugC 7
4	DrugA	3	14 DrugB 4	24 DrugC 5
5	DrugA	2	15 DrugB 6	25 DrugC 6
6	DrugA	4	16 DrugB 5	26 DrugC 5
7	DrugA	3	17 DrugB 8	27 DrugC 5

- ▶ Tidy:
 - ▶ One column per variable (2 columns)
 - ▶ One row per observation (27 rows)
- ▶ “Long format”.
- ▶ Analysis eg. `aov(pain~drug, data=migraine2)`.

Side-by-side boxplots

```
> boxplot(pain~drug,data=migraine2)
```



spread: the inverse of gather

- ▶ Our 27 migraine sufferers only tested one drug each (“between-subjects design”).
- ▶ But compare this (scores for 4 subjects in an ADHD study):

Subject	When	Score			
			3	t1	20
1	t1	13	3	t2	13
1	t2	10	3	t3	17
1	t3	12	4	t1	26
2	t1	30	4	t2	17
2	t2	20	4	t3	20
2	t3	26			

- ▶ Tidy-ish (one variable per column).
- ▶ But each subject split over 3 rows.
- ▶ “Within-subject design”: repeated-measures analysis requires measurements for same subject in same line.
- ▶ Need to “undo” gather: spread.

Spreading these data

- ▶ When column is names of variables:

```
> adhd=read.table("adhd.txt",header=T)  
> spread(adhd,When,Score)
```

	Subject	t1	t2	t3
1	1	13	10	12
2	2	30	20	26
3	3	20	13	17
4	4	26	17	20

- ▶ Can turn this back into long format with `gather`.

Columns containing combined information

- ▶ Common way to display data is *contingency table*:

Species	Disease present		Disease absent	
	Location X	Location Y	Location X	Location Y
A	44	12	38	10
B	28	22	20	18

- ▶ File might be formatted like this:

Species	px	py	ax	ay
A	44	12	38	10
B	28	22	20	18

- ▶ Columns actually encode *two* variables: whether or not disease is present, and the location.

Tidying this table (1)

- Gather up columns that are all frequencies:

```
> disease=read.table("disease.txt",header=T)
> tmp=gather(disease,dis.loc,frequency,-Species)
> tmp
```

	Species	dis.loc	frequency
1	A	px	44
2	B	px	28
3	A	py	12
4	B	py	22
5	A	ax	38
6	B	ax	20
7	A	ay	10
8	B	ay	18

Tidying this table (2)

- ▶ Column now called `dis.loc` contains *two* variables: presence or absence of disease and location.
- ▶ `separate`:

```
> separate(tmp, dis.loc, c("disease", "location"), 1)
```

	Species	disease	location	frequency
1	A	p	x	44
2	B	p	x	28
3	A	p	y	12
4	B	p	y	22
5	A	a	x	38
6	B	a	x	20
7	A	a	y	10
8	B	a	y	18

- ▶ Data frame, variable to split, what to split into, split after character 1.

dplyr: general data manipulation

- ▶ selecting rows from a data frame by value
- ▶ selecting columns from a data frame by name
- ▶ creating new variables from old ones
- ▶ summarizing variables, possibly by groups

Use (long) ADHD study results in `adhd` for example.

The ADHD data

```
> adhd
```

	Subject	When	Score
1	1	t1	13
2	1	t2	10
3	1	t3	12
4	2	t1	30
5	2	t2	20
6	2	t3	26
7	3	t1	20
8	3	t2	13
9	3	t3	17
10	4	t1	26
11	4	t2	17
12	4	t3	20

```
> library(dplyr)
```


Selecting rows (1)

- ▶ The rows for subject 2:

```
> filter(adhd, Subject==2)
```

	Subject	When	Score
1	2	t1	30
2	2	t2	20
3	2	t3	26

- ▶ The time-3 scores:

```
> filter(adhd, When=="t3")
```

	Subject	When	Score
1	1	t3	12
2	2	t3	26
3	3	t3	17
4	4	t3	20

Selecting rows (2)

- ▶ The scores bigger than 25:

```
> filter(adhd, Score>25)
```

	Subject	When	Score
1	2	t1	30
2	2	t3	26
3	4	t1	26

- ▶ Scores either for subject 2 or score 25+:

```
> filter(adhd, Subject==2 | Score>25)
```

	Subject	When	Score
1	2	t1	30
2	2	t2	20
3	2	t3	26
4	4	t1	26

Selecting columns

Name the ones you want to keep or to omit. Thus:

```
> select(adhd, c(Subject,  
+   Score))
```

	Subject	Score
1	1	13
2	1	10
3	1	12
4	2	30
5	2	20
6	2	26
7	3	20
8	3	13
9	3	17
10	4	26
11	4	17
12	4	20

```
> select(adhd, -When)
```

	Subject	Score
1	1	13
2	1	10
3	1	12
4	2	30
5	2	20
6	2	26
7	3	20
8	3	13
9	3	17
10	4	26
11	4	17
12	4	20

Creating new variables

- Score was out of 30. Turn into a percentage:

```
> tmp=mutate(adhd,pct=Score/30*100)
```

```
> head(tmp)
```

	Subject	When	Score	pct
1	1	t1	13	43.33333
2	1	t2	10	33.33333
3	1	t3	12	40.00000
4	2	t1	30	100.00000
5	2	t2	20	66.66667
6	2	t3	26	86.66667

Create percent and get rid of old score

```
> tmp=mutate(adhd,pct=Score/30*100)
> tmp2=select(tmp,-Score)
> head(tmp2)
```

	Subject	When	pct
1	1	t1	43.33333
2	1	t2	33.33333
3	1	t3	40.00000
4	2	t1	100.00000
5	2	t2	66.66667
6	2	t3	86.66667

- Created a lot of temporary variables. Can we do better?

The “pipe” operator

Same effect as previous, but “use output from last step as input to next one”:


```
> adhd %>%  
+   mutate(pct=Score/30*100) %>%  
+   select(-Score) %>%  
+   head()
```

	Subject	When	pct
1	1	t1	43.33333
2	1	t2	33.33333
3	1	t3	40.00000
4	2	t1	100.00000
5	2	t2	66.66667
6	2	t3	86.66667

In a pipe, the initial data frame argument to any function disappears.

When the data frame isn't first

```
> adhd %>%  
+   mutate(pct=Score/30*100) %>%  
+   boxplot(pct~Subject,data=.)
```



Saving pipe output to a variable

I like this way:

```
> adhd %>%  
+   mutate(pct=Score/30*100) %>%  
+   select(-Score) -> adhd.2  
> head(adhd.2)
```

	Subject	When	pct
1	1	t1	43.33333
2	1	t2	33.33333
3	1	t3	40.00000
4	2	t1	100.00000
5	2	t2	66.66667
6	2	t3	86.66667

The “right-assignment” saves the data frame on the left (the output from `select`) into the variable on the *right*.

Summarizing variables

- ▶ Mean pct, well, duh:

```
> summarize(adhd.2, mean=mean(pct))
```

```
mean
```

```
1 62.22222
```

- ▶ Summaries by *groups* takes an extra step:

```
> adhd.2 %>%
```

```
+   group_by(Subject) %>%
```

```
+   summarize( pct.mean=mean(pct),
```

```
+               pct.sd=sd(pct) )
```

```
Source: local data frame [4 x 3]
```

	Subject (int)	pct.mean (dbl)	pct.sd (dbl)
1	1	38.88889	5.091751
2	2	84.44444	16.777410
3	3	55.55556	11.706282
4	4	70.00000	15.275252

dplyr and SQL


dplyr	SQL
select	select
filter	where
mutate	alter table
summarize	select mean(x) ...

ggplot: a grammar of graphics

- ▶ Base R graphics are functional, once you get used to the quirks.
- ▶ I learned base graphics by seeing a lot of examples.
- ▶ Hadley Wickham used a “grammar of graphics”, implemented in ggplot2.
- ▶ Takes some getting used to, but once you do, *everything is consistent*.
- ▶ Separates:
 - ▶ what to plot
 - ▶ how to plot it
- ▶ Layers on plot constructed by *adding* (literally).


Histogram of ADHD percents

```
> library(ggplot2)
> p=ggplot(adhd.2,aes(x=pct))
> p+geom_histogram()
```




A better histogram

```
> p=ggplot(adhd.2,aes(x=pct))  
> p+geom_histogram(binwidth=10)
```




Histogram with density curve

```
> p=ggplot(adhd.2,aes(x=pct))  
> p+geom_histogram(binwidth=10,aes(y=..density..))+  
+   geom_density(col="blue")
```




Boxplot of percents by When

```
> p=ggplot(adhd.2,aes(x=When,y=pct))  
> p+geom_boxplot()
```




Boxplot of percents by subject

```
> p=ggplot(adhd.2,aes(x=factor(Subject),y=pct))  
> p+geom_boxplot()
```




Line plot of subject's percents by When

```
> p=ggplot(adhd.2,aes(x=When,y=pct,group=Subject))  
> p+geom_line()
```



Same line plot, with colours

```
> p=ggplot(adhd.2,aes(x=When,y=pct,group=Subject))  
> p+geom_line(aes(colour=factor(Subject)))
```




dplyr and ggplot2

- ▶ The first argument to `ggplot` is a data frame, so `ggplot` plays nicely with pipes.
- ▶ For example, to start from `adhd`, produce the percentages, get rid of the original scores, and then make the line plot, we can do all this:

```
> adhd %>%  
+   mutate(pct=Score/30*100) %>%  
+   select(-Score) %>%  
+   ggplot(aes(x=When,y=pct,group=Subject)) +  
+     geom_line(aes(colour=factor(Subject)))
```

Output



That contingency table

Species	Disease present		Disease absent	
	Location X	Location Y	Location X	Location Y
A	44	12	38	10
B	28	22	20	18


- ▶ Make a plot of disease present/absent by location and species.
- ▶ Technique in ggplot called **faceting**.
- ▶ Code we had before, rewritten with pipes:

```
> disease %>% gather(dis.loc,frequency,-Species) %>%  
+   separate(dis.loc,c("disease","location"),1) -> dis.2  
> head(dis.2,4)
```

	Species	disease	location	frequency
1	A	p	x	44
2	B	p	x	28
3	A	p	y	12
4	B	p	y	22


Without considering species and location

```
> dis.2 %>%  
+   ggplot(aes(x=disease,weight=frequency)) +  
+   geom_bar()
```




By species and location

```
> dis.2 %>%  
+   ggplot(aes(x=disease,weight=frequency)) +  
+   geom_bar() + facet_grid(Species ~ location)
```




Stacked bars, with black magic

```
> dis.2 %>%  
+   ggplot(aes(x=factor(1),weight=frequency,fill=disease))  
+   geom_bar(aes(y=(..count..)/  
+   tapply(..count..,..PANEL..,sum)[..PANEL..])) +  
+   facet_grid(Species~location)
```



Or as pie charts, if you must

```
> dis.2 %>%  
+   ggplot(aes(x=factor(1),weight=frequency,fill=disease)) +  
+   geom_bar(aes(y=(..count..)/  
+   tapply(..count..,..PANEL..,sum)[..PANEL..])) +  
+   facet_grid(Species~location) +  
+   coord_polar(theta="y")
```



Soccer goalscoring

- The English Premier league this year:

```
> premier=read.csv("premier.csv",header=T)
> tail(premier)
```

	date	name	name.1	score
104	2015-10-31	West Bromwich Albion	Leicester City	2 - 3
105	2015-10-31	Crystal Palace	Manchester United	0 - 0
106	2015-10-31	Watford	West Ham United	2 - 0
107	2015-10-31	Swansea City	Arsenal	0 - 3
108	2015-11-01	Southampton	AFC Bournemouth	2 - 0
109	2015-11-01	Everton	Sunderland	6 - 2


Sorting out those goals

```
> premier %>%  
+   separate(score, c("home", "away"), " - ") -> goals  
> tail(goals)
```

	date	name	name.1	home	away
104	2015-10-31	West Bromwich Albion	Leicester City	2	3
105	2015-10-31	Crystal Palace	Manchester United	0	0
106	2015-10-31	Watford	West Ham United	2	0
107	2015-10-31	Swansea City	Arsenal	0	3
108	2015-11-01	Southampton	AFC Bournemouth	2	0
109	2015-11-01	Everton	Sunderland	6	2


Scatterplot of goals

```
> goals %>%  
+   ggplot(aes(x=home,y=away)) + geom_point()
```



Jittering to solve overplotting

```
> goals %>% ggplot(aes(x=home, y=away)) + geom_jitter()
```



Comments


- ▶ Default jitter is “40% of resolution of data” in each direction, here, 0.4 goals.
- ▶ Can colour each point by the month in which the game was played.
- ▶ Alternative is to plot at actual points, but have size of symbol proportional to frequency.

Jittered plot coloured by month

- ▶ Uses another Wickham package `lubridate` to extract the month of a date.
- ▶ My dates are text strings, so convert to R Dates first.
- ▶ Add the colour to the `aes` for the `ggplot`.


```
> library(lubridate)
> goals %>%
+   mutate(thedate=as.Date(date)) %>%
+   mutate(themonth=month(thedate,label=T)) %>%
+   ggplot(aes(x=home,y=away,colour=themonth)) +
+     geom_jitter()
```

The plot



Original points, plotted according to frequency

```
> goals %>% ggplot(aes(x=home,y=away)) +  
+   stat_sum(aes(size = factor(..n..)), geom = "point")
```



Thank you

for your attention!

<http://www.utsc.utoronto.ca/~butler/rug/>